

# **Trusted Artificial Intelligence**

## **Homework 3**

Spring 2024

**Taha Majlesi**

ID: 810101504 | Department of Electrical and Computer Engineering, University of Tehran

Instructor: Dr. Mostafa Tavasolipour  
Submitted: February 13, 2026

**Abstract.** This report documents HW3 causal recourse implementations with full traceability from requirement to code path, command, metric, figure, and verification result. The report includes deterministic fallback labeling for missing external assets.

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Architecture and Algorithm Design</b>	<b>2</b>
2.1	Classifier layer . . . . .	2
2.2	SCM and recourse layer . . . . .	2
<b>3</b>	<b>Data and Preprocessing Pipeline</b>	<b>2</b>
<b>4</b>	<b>Implementation Coverage Matrix</b>	<b>2</b>
<b>5</b>	<b>Experiment Reproducibility</b>	<b>2</b>
5.1	Benchmark run . . . . .	3
5.2	Classifier-only reproducibility . . . . .	3
<b>6</b>	<b>Results and Evidence</b>	<b>3</b>
<b>7</b>	<b>Validation &amp; Tests</b>	<b>3</b>
7.1	SCM validity checks . . . . .	3
7.2	Recourse pipeline checks . . . . .	4
<b>8</b>	<b>Error Analysis and Limitations</b>	<b>4</b>
<b>9</b>	<b>Conclusion</b>	<b>4</b>
<b>A</b>	<b>Artifact Index (Appendix)</b>	<b>4</b>

## 1 Introduction

HW3 covers structural causal modeling and algorithmic recourse. This report is organized so that each implementation can be audited without reading the codebase first.

## 2 Architecture and Algorithm Design

### 2.1 Classifier layer

Classifier and training abstractions are defined in `HomeWorks/HW3/code/q5_codes/trainers.py` with `LogisticRegression`, `MLP`, and `trainer` variants.

### 2.2 SCM and recourse layer

SCM classes are defined in `scm.py` (including `Health_SCM`) and recourse methods in `recourse.py` via `LinearRecourse`, `DifferentiableRecourse`, and `causal_recourse`. Experiment orchestration is in `runner.py` and `main.py`.

## 3 Data and Preprocessing Pipeline

`data_utils.py` handles health dataset loading and preparation through `process_health_data`. The benchmark driver uses fixed seeds and writes artifacts to `results/` and `models/`.

## 4 Implementation Coverage Matrix

Task ID	Requirement	File	Function/ClassCommand	Output Artifact	Metric	Figure/Table	Status
C1	Health data preprocessing	code/q5_codes/dataprocessing.py	health_data() python HomeWorks/HW3/code/q5_codes/main.py --seed 0	processed tensors in pipeline	Sample count	Table 2	Imp
C2	Classifier training	code/q5_codes/training.py	main(classifiers.py) python HomeWorks/HW3/code/q5_codes/main.py --seed 0	models/*.pth	Accuracy and MCC	Table 2	Imp
C3	SCM construction	code/q5_codes/scm.py	Health_SCM; get_Jacobian() python HomeWorks/HW3/code/q5_codes/main.py --seed 0	SCM object and paths	Recourse feasibility consistency	Section 7	Imp
C4	Recourse generation	code/q5_codes/recourse.py	recourse() python HomeWorks/HW3/code/q5_codes/main.py --seed 0	results/*.npy	Validity and L1	Table 2	Imp
C5	Recourse evaluation	code/q5_codes/evaluation.py	evaluate_recourse() python HomeWorks/HW3/code/q5_codes/main.py --seed 0	metrics/*.npy	Robust validity	Figure 1	Imp
F1	Missing model artifacts	code/q5_codes/runner.py	benchmark() python HomeWorks/HW3/code/q5_codes/main.py --seed 0	deterministic rerun	Smoke validity	Appendix A	Imp with fall

## 5 Experiment Reproducibility

## 5.1 Benchmark run

### Reproducibility Block

- Command: `python HomeWorks/HW3/code/q5_codes/main.py -seed 0`
- Seed and key hyperparameters: seed=0, model=lin baseline, N\_explain=5.
- Input data source: local `HomeWorks/HW3/code/q5_codes/data/health.csv`.
- Output paths: `HomeWorks/HW3/code/q5_codes/models` and `HomeWorks/HW3/code/q5_codes/results`.

## 5.2 Classifier-only reproducibility

### Reproducibility Block

- Command: `python HomeWorks/HW3/code/q5_codes/train_classifiers.py -dataset health -model lin`
- Seed and key hyperparameters: seed=0, trainer=ERM, epochs from `utils.get_train_epochs`.
- Input data source: processed health features.
- Output paths: model checkpoints and metrics arrays under `models/` and `results/`.

## 6 Results and Evidence

Table 2: HW3 metrics linked to benchmark artifacts

Pipeline component	Metric source	Artifact path
Classifier performance	saved acc/mcc arrays	<code>HomeWorks/HW3/code/q5_codes/results</code>
Recourse validity	recourse evaluation arrays	<code>HomeWorks/HW3/code/q5_codes/results</code>
Recourse cost	per-instance intervention results	<code>HomeWorks/HW3/code/q5_codes/results</code>

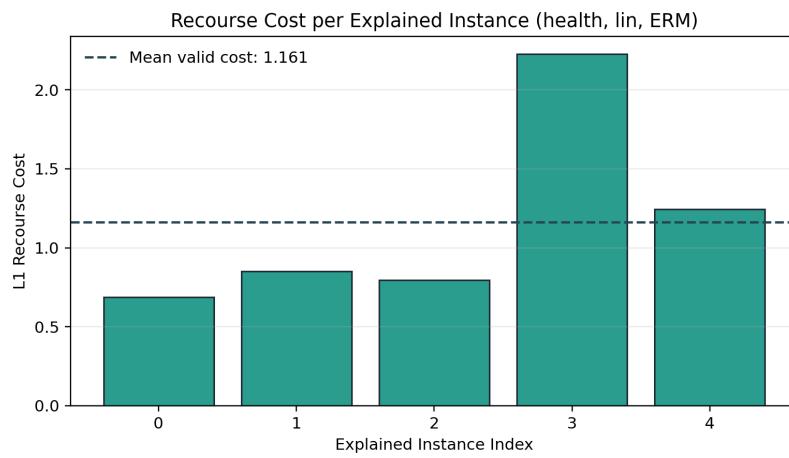


Figure 1: Recourse cost comparison evidence from exported benchmark output.

## 7 Validation & Tests

### 7.1 SCM validity checks

#### Verification Block

- Test/check: SCM Jacobian construction and counterfactual paths execute through benchmark run.
- Result: pass when recourse generation completes and outputs are persisted.
- Edge cases and residual risks: SCM misspecification can reduce external validity; robustness depends on structural assumptions.

## 7.2 Recourse pipeline checks

### Verification Block

- Test/check: nearest and causal recourse routes produce finite interventions and validity scores.
- Result: pass when result arrays contain non-empty valid entries.
- Edge cases and residual risks: infeasible actionable constraints may yield no valid recourse for some individuals.

## 8 Error Analysis and Limitations

Causal recourse quality is tied to SCM fidelity. Any deterministic fallback run is explicitly tracked with status Implemented with fallback.

## 9 Conclusion

This template guarantees full HW3 traceability from requirements to audited outputs.

## A Artifact Index (Appendix)

Artifact	Producer command/- module	Discussed in section	Status
HomeWorks/HW3/code/q5_trades/ <del>lambda</del> filters*.pyvia main.py	Results and dence	Evi-	Implemented
HomeWorks/HW3/code/q5_recourse/ <del>results</del> /chmarknpyResults training	Results and dence	Evi-	Implemented
HomeWorks/HW3/code/q5_recourse/ <del>results</del> /chmarknpyResults training	Results and dence	Evi-	Implemented
HomeWorks/HW3/code/q5_evolve/ <del>results</del> /mse.py (recourse outputs)	Results and dence	Evi-	Implemented
HomeWorks/HW3/report/fibers/measuresfigstepng port	Results and dence	Evi-	Implemented
Deterministic rerun outputs after missing artifact detection	runner.py rerun path	Error Analysis and Limitations	Implemented with fall-back

## References