

Complete Causal Recourse Implementation on Health Data

(IEEE-Style Report for Trusted AI HW3, Question 5)

Taha Majlesi, Student ID 810101504

Department of Electrical and Computer Engineering, University of Tehran

Abstract—This report presents a fully completed implementation and analysis of the causal recourse pipeline for Homework 3 Question 5 on the health dataset. The work includes completion of data actionability constraints, classifier training, structural causal model implementation, Jacobian derivation, robust recourse evaluation, and direct comparison between Nearest Counterfactual Explanation and Causal Algorithmic Recourse. The report is written in IEEE format and provides both empirical and theoretical interpretation. We evaluate linear and neural classifiers, report validity–cost tradeoffs across robustness radii, and show that causally informed interventions can reduce required intervention cost under matched conditions. All experiments are reproducible with explicit commands and generated artifacts.

Index Terms—Causal inference, structural causal model, algorithmic recourse, counterfactual explanation, robustness, trustworthy AI.

I. INTRODUCTION

Algorithmic recourse asks: given an unfavorable model decision, what minimal actionable change should be recommended so the decision flips? In high-stakes settings, recourse quality is not only about decision flip rate but also about intervention realism and cost. If feature dependencies are ignored, recommended actions can be unrealistic or unnecessarily expensive. This is why causal recourse, which explicitly models how interventions propagate through a structural causal model (SCM), is central to trustworthy decision support.

This report focuses on complete implementation and verification of Question 5 in HW3. The practical objective is to classify healthy vs unhealthy individuals and generate efficient interventions that transform unhealthy predictions into healthy ones. Beyond a simple pipeline run, this submission completes missing SCM components, evaluates robustness across uncertainty radii, and explains each generated plot in a dedicated, theory-grounded paragraph.

II. THEORETICAL BACKGROUND

A. Counterfactual and Causal Recourse

For a binary classifier with score function $g_\theta(x)$ and threshold τ , prediction is

$$\hat{y} = \mathbb{I}[\sigma(g_\theta(x)) \geq \tau]. \quad (1)$$

Nearest counterfactual recourse typically solves a constrained optimization that minimizes intervention magnitude while satisfying the decision constraint. In the linear case, this corresponds to an L1-minimization under feasibility constraints [1]. Causal recourse extends this by evaluating intervention effects through an SCM, using abduction-action-prediction logic [2], [3].

B. Robust Linear Recourse Geometry

Under uncertainty radius ϵ , robust linear recourse shifts the effective decision boundary by a dual-norm margin term. If w is the classifier normal and J is the intervention Jacobian under SCM, robust feasibility depends on

$$\langle w, x + Ja \rangle \geq b + \|J^\top w\|_2 \epsilon. \quad (2)$$

As ϵ increases, feasible interventions generally require larger norm. Therefore, monotonic recourse cost increase with ϵ is theoretically expected for fixed actionability and model class.

C. Differentiable Recourse for Nonlinear Models

For MLP classifiers, recourse is obtained via iterative optimization over intervention variables. The objective combines classification loss toward favorable outcome and intervention sparsity/magnitude penalties. Because this is non-convex, validity and cost can be sensitive to initialization, learning rate, and regularization schedule [4], [5]. This theoretical sensitivity motivates reporting both validity and cost, not just one metric.

III. IMPLEMENTATION COMPLETION FOR Q5

A. Q5.1 Data Processing and Actionability

In `code/q5_codes/data_utils.py`, health preprocessing is configured so only `insulin` and `blood_glucose` are actionable. Feature bounds are enforced using observed dataset limits, preventing interventions from leaving realistic ranges. Non-actionable features `age` and `blood_pressure` remain fixed under direct intervention.

B. Q5.2 Running on 10 Unhealthy Individuals

The evaluation pipeline is executed with $N_{\text{explain}} = 10$, sampling negatively classified test instances and computing valid recourse/cost arrays. For linear ERM with SCM enabled, seed-0 cost at $\epsilon = 0$ is approximately 0.909, and the multi-seed mean is 0.889.

C. Q5.3 and Q5.4 Completing Health_SCM and Jacobian

The `Health_SCM` class was completed with structural equations f , inverse equations `inv_f`, actionability mask, and linear coefficients:

$$X_1 = U_1, \quad (3)$$

$$X_2 = \frac{1}{18}X_1 + U_2, \quad (4)$$

$$X_3 = 2.0X_1 + 1.05X_2 + U_3, \quad (5)$$

$$X_4 = 0.4X_2 + 0.3X_3 + U_4. \quad (6)$$

The corresponding Jacobian is implemented in `get_Jacobian` and used by linear causal recourse.

D. Q5.5 and Q5.6 SCM-On Rerun and Method Comparison

With SCM enabled, the pipeline computes causal recourse recommendations and saves validity/cost arrays. Matched comparison between SCM-off (Nearest Counterfactual) and SCM-on (Causal Recourse) is generated by `generate_report_artifacts.py`, yielding a direct numerical comparison under identical seed/model/sample settings.

IV. COMPLETE CODE WALKTHROUGH

A. End-to-End Control Flow

The executable entry point is `code/q5_codes/main.py`. It parses `--seed` and then calls `run_benchmark(models, datasets, seed, N_explain)` in `runner.py`. Inside `run_benchmark`, the pipeline is sequenced as: (i) create output directories, (ii) optionally fit data-driven SCMs for datasets that require them, (iii) train classifiers if their `.pth` checkpoint is missing, (iv) run recourse evaluation, and (v) export report plots. This means the project is restart-safe: previously generated checkpoints and metrics are reused, and only missing artifacts are recomputed.

B. Data Layer (`data_utils.py`)

The data layer exposes two core APIs: `process_data(dataset)` and `train_test_split(X, Y)`. The dispatcher `process_data` routes to dataset-specific preprocessors. For HW3-Q5, `process_health_data()` loads `health.csv`, extracts the four modeled variables (`age`, `insulin`, `blood_glucose`, `blood_pressure`), standardizes them using `StandardScaler`, and returns a constraints dictionary with actionable indices, monotonic direction constraints, and per-feature intervention limits in standardized space. The important implementation detail is

that feature bounds are computed from raw min/max and then mapped into normalized coordinates; this keeps optimization numerically stable while still enforcing physically meaningful limits.

C. Model Layer (`trainers.py` and `train_classifiers.py`)

Model construction and optimization are separated. `train_classifiers.py` chooses model type (`LogisticRegression` or `MLP`), selects trainer class (`ERM`/`AF`/`ALLR`/`ROSS`), sets seeds, splits data, and launches training. In `trainers.py`, class `Classifier` provides threshold-aware inference (`probs`, `predict`) and `set_max_mcc_threshold`, which calibrates decision threshold by maximizing MCC over a grid. `LogisticRegression.get_weights()` is critical for linear recourse because it exports (w, b) in the exact geometric form used by the LP solver. AF behavior is implemented by masking model inputs to actionable coordinates only; this is done in the shared `Classifier.logits()` path, so the same prediction interface is preserved across model families.

D. SCM Layer (`scm.py`)

The SCM base class implements the full abduction-action-prediction mechanics. `Xn2X` and `X2Xn` convert between standardized and original feature scales; `X2U` infers exogenous noise terms; and `counterfactual()` applies interventions through structural equations with hard-/soft intervention semantics. The completed `Health_SCM` defines forward equations `self.f`, inverse equations `self.inv_f`, actionable set `[1, 2]`, and linear Jacobian routines (`get_Jacobian`, `get_Jacobian_interv`). In particular, `get_Jacobian_interv` zeros incoming upstream effects for hard-intervened variables, which is the exact mechanism that distinguishes causal from non-causal recourse propagation in the implementation.

E. Recourse Solver Layer (`recourse.py`)

This file contains both linear and nonlinear recourse engines. `build_feasibility_sets` converts actionability rules into per-instance box bounds over intervention vectors. `LinearRecourse.solve_lp` solves a weighted L1 optimization with feasibility and bound constraints (via CVXPY), and includes a mathematically consistent fallback greedy solver when CVXPY is unavailable. `DifferentiableRecourse.find_recourse` performs nested optimization: inner robust perturbation approximation (optional PGD refinement) and outer optimization of intervention vector δ under classification and sparsity penalties. Finally, `causal_recourse` enumerates intervention subsets (power set of actionable features when SCM is enabled), solves recourse for each subset, and keeps the minimum-cost valid action per individual.

F. Evaluation Layer (`evaluate_recourse.py`)

Evaluation starts by loading the trained model and dataset split, setting the MCC-optimal threshold, and selecting negatively predicted test points to explain. The linear branch computes robust threshold shift using $\|J^\top w\|_2 \epsilon$, then runs LP-based recourse; the MLP branch uses differentiable recourse with hyperparameters from `utils.get_recourse_hyperparams`. Results are saved in a deterministic naming scheme (`_ids.npy`, `_valid.npy`, `_cost.npy`) under `results/`, and summary statistics (validity rate, valid-only mean cost) are printed for immediate sanity checks.

G. Reporting Layer (`generate_report_artifacts.py` and `plot_report_figures.py`)

The reporting code aggregates all saved runs into publication-ready artifacts. `generate_report_artifacts.py` parses model filenames, reloads models, recomputes classifier metrics consistently, merges them with recourse outputs for each (model, trainer, ϵ , seed), writes machine-readable CSV summaries, and renders final figures used in the report. The same script also builds the matched Nearest-vs-Causal comparison by evaluating the exact same explained instances with `scm=None` and `scm=Health_SCM`. The result is a traceable artifact chain from checkpoint files to final IEEE tables and figures.

H. Utility and Naming Conventions (`utils.py`)

`utils.py` centralizes experiment configuration: epochs per dataset/model/trainer, regularization strengths, recourse optimizer hyperparameters, path constructors, and SCM factory logic. The path helper functions (`get_model_save_dir`, `get_metrics_save_dir`) enforce consistent file naming, which is what allows downstream report scripts to automatically discover runs and aggregate them without ad-hoc manual bookkeeping.

I. Implementation Correctness Summary

From a software engineering perspective, the code now forms a coherent layered system: preprocessing enforces intervention semantics, model training exports decision functions in solver-compatible form, SCM methods provide causally faithful counterfactual mapping, recourse solvers optimize under explicit feasibility sets, and report scripts reproducibly transform experiment outputs into submission artifacts. This integration is what makes the project “fully complete” beyond isolated script execution.

V. EXPERIMENTAL PROTOCOL

A. Environment and Reproducibility

All runs use:

- Python environment: `/Users/tahamajs/Documents/`
- Code root: `HomeWorks/HW3/code/q5_codes`
- Report root: `HomeWorks/HW3/report`

TABLE I
MODEL AND REOURSE SETTINGS USED IN THIS REPORT

Configuration	Seeds	ϵ set	N_{explain}
lin-ERM	0,1,2	{0.0, 0.1, 0.2}	10
lin-AF	0,1,2	{0.0, 0.1, 0.2}	10
mlp-ERM	0,1,2	{0.0, 0.1, 0.2}	10
mlp-AF	0,1,2	{0.0, 0.1, 0.2}	10

TABLE II
CLASSIFIER QUALITY (MEAN \pm STD ACROSS AVAILABLE SEEDS)

Configuration	Accuracy	MCC
lin-ERM	0.900 ± 0.001	0.805 ± 0.001
lin-AF	0.899 ± 0.002	0.798 ± 0.005
mlp-ERM	0.997 ± 0.002	0.995 ± 0.003
mlp-AF	0.998 ± 0.001	0.995 ± 0.002




Fig. 1. Classifier metrics by model/trainer.

B. Evaluated Configurations

C. Generated Analysis Artifacts

The script `generate_report_artifacts.py` produces:

- `results/health_report_summary.csv`
- `results/health_report_aggregate.csv`
- `results/nearest_vs_causal_lin_seed0.csv`
- `results/health_instance_costs.csv`
- `results/health_action_profiles.csv`
- Plot files under `report/figures/`

VI. RESULTS AND COMPLETE PLOT EXPLANATIONS

A. Classifier Performance Summary

Complete interpretation of Fig. 1: This plot shows two clear regimes: linear models (ERM and AF) have similar predictive strength around 0.899–0.900 accuracy and 0.798–0.805 MCC, while MLP models (ERM and AF) are substantially higher near 0.997–0.998 accuracy and about 0.995 MCC. Theoretically, this supports the claim that actionability masking does not impose a major predictive penalty when actionable variables already capture most task-relevant signal. At the same time, the figure emphasizes a key recourse principle: predictive quality and intervention quality are different objectives. Even when discrimination is excellent, intervention feasibility and cost depend on the geometry of actionable directions, the

TABLE III
RECOURSE OUTCOMES (MEAN ACROSS SEEDS)

Configuration	ϵ	Valid rate	Mean valid cost
lin-ERM	0.0	1.000	0.889
lin-ERM	0.1	1.000	1.004
lin-ERM	0.2	1.000	1.120
lin-AF	0.0	1.000	0.701
lin-AF	0.1	1.000	0.823
lin-AF	0.2	1.000	0.946
mlp-ERM	0.0	0.867	1.177
mlp-ERM	0.1	0.900	1.334
mlp-ERM	0.2	0.900	1.150
mlp-AF	0.0	0.967	1.793
mlp-AF	0.1	0.967	1.971
mlp-AF	0.2	0.933	1.988




Fig. 2. Valid recourse rate vs robustness radius ϵ .

causal Jacobian, and the optimization dynamics used to find recourse.

B. Validity–Cost Tradeoff Across Robustness Radius

Complete interpretation of Fig. 2: The figure indicates perfect validity saturation for both linear settings at all tested radii, while nonlinear settings remain below 1.0 with model-dependent behavior (MLP-AF above MLP-ERM but not perfect). This pattern is theoretically consistent with convex versus non-convex recourse search: linear robust recourse has explicit Jacobian-shifted constraints and a stable feasible-set characterization, whereas MLP recourse is obtained by iterative gradient steps over a non-convex objective and can terminate in local basins or near-boundary states that do not cross the threshold. The higher MLP-AF validity here suggests that constraining classifier dependence to actionable coordinates can improve optimization alignment, yet finite-step optimization and heterogeneous instance geometry still prevent guaranteed validity.

Complete interpretation of Fig. 3: For both linear models, intervention cost increases nearly linearly with ϵ , which directly matches robust optimization theory: larger uncertainty requires a larger worst-case margin, hence larger minimum L1 action. AF remains strictly cheaper than ERM in the linear case, supporting the geometric view that actionable masking can rotate effective decision sensitivity toward feasible intervention directions. In nonlinear settings, costs are markedly higher and more variable, and MLP-AF is especially expensive despite higher validity. This is theoretically plausible because




Fig. 3. Mean valid recourse cost vs robustness radius ϵ .




Fig. 4. Per-instance recourse costs for explained unhealthy individuals.

gradient-based search may find valid but distant interventions when loss curvature, step-size schedule, and action-penalty coupling favor large moves in a subset of hard instances.

C. Instance-Level Cost Distribution

Complete interpretation of Fig. 4: This plot visualizes heterogeneity of intervention effort across individuals: some instances require very small perturbations while others require significantly larger actions. Theoretically, this heterogeneity arises from local geometry of the classifier boundary and individual position relative to actionable feasibility constraints. Points near the boundary and aligned with high-gain actionable directions need small interventions; points deeper in the unfavorable region, or constrained by directional/box bounds, require larger L1 actions. Therefore, average recourse cost should always be interpreted together with distributional spread, not as a single universal burden.

D. Nearest Counterfactual vs Causal Recourse

Complete interpretation of Fig. 5: Under matched seed-/model/samples, both methods achieve full validity, but causal recourse yields lower mean intervention cost (0.589 versus 0.733). Theoretically, SCM-aware optimization can leverage causal amplification: modifying an actionable parent induces beneficial downstream shifts through structural equations, increasing classifier score per unit direct intervention. In contrast, nearest counterfactual search without SCM treats correlated descendants as independent dimensions and may spend




Fig. 5. Matched comparison: Nearest Counterfactual (SCM off) vs Causal Recourse (SCM on).




Fig. 6. Validity-cost frontier across model/trainer/epsilon settings.

action budget redundantly. This cost gap therefore reflects an efficiency benefit from structural knowledge, not merely a random optimization artifact, and aligns with intervention-based recourse theory.

E. Expanded Diagnostic Features for Complete Understanding

Complete interpretation of Fig. ??: This frontier plot makes explicit that recourse quality is a multi-objective operating point rather than a single score. Points near the top-left are preferable (high validity, low cost), while downward or rightward shifts indicate weaker practical recourse quality. The linear AF family sits on a favorable region with both perfect validity and lower cost than linear ERM, while nonlinear settings occupy higher-cost regions despite strong classifier accuracy. Theoretically, this figure is useful because it separates predictive performance from intervention burden and visualizes the Pareto-like tradeoff that must be reported for trustworthy deployment.

Complete interpretation of Fig. ??: Unlike mean-only summaries, this boxplot reveals distributional behavior and tail risk. Linear configurations show tighter spread and predictable median shifts with ϵ , indicating stable geometry under robust margin increases. Nonlinear configurations exhibit wider dispersion and heavier upper tails, implying that a subset of individuals pays substantially larger intervention cost even when average validity is acceptable. This is theoretically




Fig. 7. Per-instance recourse cost distribution by configuration and epsilon.




Fig. 8. Feature-wise mean absolute intervention magnitude (valid recourse only).




Fig. 9. Feature intervention activation rate among valid recourse actions.

important because fairness and usability concerns are often driven by high-cost tails, not by central tendency alone.

Complete interpretation of Fig. ??: This diagnostic quantifies where intervention budget is actually spent. Since only insulin and blood glucose are actionable, large action mass should concentrate on those coordinates while non-actionable dimensions remain near zero. The plotted pattern confirms this implementation behavior and also reveals model-dependent preference among actionable features, which reflects how each classifier's local gradient and SCM propagation jointly determine the most efficient direction. This offers direct interpretability: the recommended changes are not only valid but also aligned with declared actionability policy.

Complete interpretation of Fig. ??: Activation rate measures how frequently each feature is used in successful interventions. A high nonzero rate for actionable variables and near-zero

rate for non-actionable variables is the expected signature of a policy-consistent recourse system. This frequency view complements magnitude view in Fig. ??: a feature can have moderate average magnitude but very high activation frequency, indicating it is a reliable “first-step” recourse coordinate. Theoretical value comes from separating sparse-but-large actions from frequent-small actions, which correspond to different behavioral recourse strategies.

VII. DISCUSSION AND THEORETICAL IMPLICATIONS

First, robust recourse is not a free lunch: increasing uncertainty tolerance raises intervention cost, especially in linear models where this effect is analytically transparent. Second, classifier architecture alone does not determine recourse practicality. The MLP results show that near-ceiling predictive metrics can coexist with high or unstable recourse costs. Third, actionability-aware training (AF) can reduce practical intervention burden in linear settings without sacrificing classifier quality, but this benefit is not guaranteed in nonlinear optimization regimes, where curvature and initialization effects can dominate.

From a causal perspective, this homework confirms a central principle: interventions should be evaluated in a structural model, not only in observational feature space. When feature dependencies are strong, SCM-enabled recommendations can be both more realistic and cheaper.

An additional implication is deployment robustness: operational recourse systems should report uncertainty bands over seeds, initialization, and optimization hyperparameters, especially for nonlinear recourse solvers. A single-point mean can hide heavy-tail intervention costs that are unacceptable in practice. Therefore, trustworthy deployment requires both average-case performance and tail-risk monitoring (e.g., quantiles of valid cost among successful recourse cases).

VIII. EXTENDED THEORETICAL ANALYSIS

A. Linear Recourse Cost Lower Bound

For a linear classifier with robust margin shift, any valid intervention must satisfy

$$\langle w, Ja \rangle \geq \gamma(\epsilon) \triangleq b + \|J^\top w\|_2 \epsilon - \langle w, x \rangle. \quad (7)$$

By Hölder duality, a coarse lower bound on L1 action is

$$\|a\|_1 \geq \frac{\gamma(\epsilon)}{\|J^\top w\|_\infty}, \quad (8)$$

when $\gamma(\epsilon) > 0$. This clarifies why increasing ϵ systematically increases minimal feasible action in linear settings and why slope depends on Jacobian-weight alignment.

B. Why AF Can Reduce Cost Without Hurting Accuracy

AF constrains model dependence to actionable coordinates. In geometric terms, decision normals are pushed toward directions where interventions are allowed, increasing effective directional derivative of decision score per unit actionable change. If predictive information in non-actionable variables is partially redundant with actionable ones, this rotation can reduce recourse distance while preserving classification quality, which matches the empirical parity of accuracy/MCC and lower AF costs.

C. Causal Amplification Mechanism

Let an intervention apply on variable set S . Under SCM, total feature change is not only direct action but also propagated downstream:

$$\Delta x_{\text{total}} = J_{SAS}. \quad (9)$$

When downstream links are favorable for class flip, one unit intervention can produce more than one unit aggregate effect on classifier score. Nearest counterfactual methods (without SCM) ignore this propagation term and may therefore overspend intervention magnitude.

D. Validity-Cost Frontier Interpretation

Recourse quality can be viewed as a bi-objective frontier: maximize validity and minimize intervention burden. Linear models in this report sit near a high-validity region with predictable cost growth as robustness tightens. MLP settings display frontier instability due optimization non-convexity; therefore, robust deployment should report confidence intervals, not single-point estimates, and include optimization diagnostics.

E. Nonlinear Recourse Curvature Effect

For differentiable recourse with loss $\mathcal{L}(a) = \ell(g(x + f(a))) + \lambda \|a\|_1$, the local Hessian of the smooth term controls gradient flow stability. In regions of high curvature, a fixed step-size can oscillate or overshoot toward higher-cost valid points. This offers a theoretical explanation for observing high validity but inflated action magnitudes in some MLP-AF runs: optimization reaches feasibility, but not low-cost local minima. In practice, line-search or adaptive trust-region updates can reduce this gap.

F. SCM Misspecification Consideration

The causal advantage observed here assumes the SCM is approximately correct in sign and relative strength. If structural coefficients are misspecified, propagated effects can be misestimated and recommended actions may become suboptimal. Nonetheless, even imperfect SCMs often provide a better inductive bias than no structure at all when domain relations are strong. This motivates future work on recourse under causal uncertainty sets, where interventions are optimized against a family of plausible SCM parameters.

IX. CONCLUSION

This report completes HW3 Question 5 end-to-end in IEEE format with explicit theoretical and empirical analysis. The software pipeline is fully runnable, missing SCM components are completed, robust evaluations are produced, and each plot is interpreted in a dedicated theory-grounded paragraph. Empirically, linear recourse is highly stable on this dataset, AF reduces intervention cost, and causal recourse outperforms nearest counterfactual in matched cost comparison while maintaining full validity.

APPENDIX A REPRODUCIBILITY COMMANDS

Listing 1. Exact commands used for the final report build

```

1 cd /Users/tahamajs/Documents/uni/truthlyAI/HomeWorks
   /HW3/code/q5_codes
2 source /Users/tahamajs/Documents/uni/venv/bin/
   activate
3
4 python main.py --seed 0
5 python generate_report_artifacts.py
6
7 cd /Users/tahamajs/Documents/uni/truthlyAI/HomeWorks
   /HW3/report
8 make pdf

```

APPENDIX B AUTO-GENERATED AGGREGATE CSV

Listing 2. Health report aggregate CSV

```

1 model,trainer,epsilon,accuracy_mean,accuracy_std,
   mcc_mean,mcc_std,valid_rate_mean,valid_rate_std,
   valid_cost_mean,valid_cost_std,runs
2 lin,AF
   ,0.0,0.8993333333333333,0.0020816659994661213,0.79
3 lin,AF
   ,0.1,0.8993333333333333,0.0020816659994661213,0.79
4 lin,AF
   ,0.2,0.8993333333333333,0.0020816659994661213,0.79
5 lin,ERM
   ,0.0,0.9003333333333333,0.0005773502691896744,0.80
6 lin,ERM
   ,0.1,0.9003333333333333,0.0005773502691896744,0.80
7 lin,ERM
   ,0.2,0.9003333333333333,0.0005773502691896744,0.80
8 mlp,AF
   ,0.0,0.9976666666666666,0.0011547005383792607,0.99
9 mlp,AF
   ,0.1,0.9976666666666666,0.0011547005383792607,0.99
10 mlp,AF
   ,0.2,0.9976666666666666,0.0011547005383792607,0.99
11 mlp,ERM
   ,0.0,0.9973333333333333,0.00152752523165193,0.9945
12 mlp,ERM
   ,0.1,0.9973333333333333,0.00152752523165193,0.9945
13 mlp,ERM
   ,0.2,0.9973333333333333,0.00152752523165193,0.9945

```

APPENDIX C AUTO-GENERATED PER-RUN CSV

Listing 3. Health report per-run summary CSV

```

1 dataset,model,trainer,seed,epsilon,accuracy,mcc,
   valid_rate,valid_cost
2 health,lin,AF
   ,0.0,0.0,0.9,0.8005852466548463,1.0,0.54514298838393

```

```

3 health,lin,AF
   ,0,0.1,0.9,0.8005852466548463,1.0,0.667209823851158
4 health,lin,AF
   ,0,0.2,0.9,0.8005852466548463,1.0,0.789276659318378
5 health,lin,AF
   ,1,0.0,0.897,0.7925895954095332,1.0,0.8750397790684868
6 health,lin,AF
   ,1,0.1,0.897,0.7925895954095332,1.0,0.9980663942579822
7 health,lin,AF
   ,1,0.2,0.897,0.7925895954095332,1.0,1.1210930094474776
8 health,lin,AF
   ,2,0.0,0.901,0.802055574103413,1.0,0.6836418661866268
9 health,lin,AF
   ,2,0.1,0.901,0.802055574103413,1.0,0.8051776933349885
10 health,lin,AF
   ,2,0.2,0.901,0.802055574103413,1.0,0.9267135204833504
11 health,mlp,AF
   ,0,0.0,0.997,0.9938787926885784,0.9,1.2796937765346632
12 health,mlp,AF
   ,0,0.1,0.997,0.9938787926885784,0.9,1.521521086494128
13 health,mlp,AF
   ,0,0.2,0.997,0.9938787926885784,0.8,1.4955620095133781
14 health,mlp,AF
   ,1,0.0,0.999,0.9979566795160255,1.0,0.8422666847705841
15 health,mlp,AF
   ,1,0.1,0.999,0.9979566795160255,1.0,0.9482423484325408
16 health,mlp,AF
   ,1,0.2,0.999,0.9979566795160255,1.0,0.9144199848175049
17 health,mlp,AF
   ,2,0.0,0.997,0.9938695652058721,1.0,3.257878613471985
18 health,mlp,AF
   ,2,0.1,0.997,0.9938695652058721,1.0,3.4426105111837386
19 health,mlp,AF
   ,2,0.2,0.997,0.9938695652058721,1.0,3.5525542467832567
20 health,lin,ERM
   ,0,0.0,0.901,0.8058258764927546,1.0,0.9083298665068884
21 health,lin,ERM
   ,0,0.1,0.901,0.8058258764927546,1.0,1.0223727404983536
22 health,lin,ERM
   ,0,0.2,0.901,0.8058258764927546,1.0,1.1364156144898185
23 health,lin,ERM
   ,1,0.0,0.9,0.8045315032435294,1.0,0.9725247265229917
24 health,lin,ERM
   ,1,0.1,0.9,0.8045315032435294,1.0,1.0892578652247575
25 health,lin,ERM
   ,1,0.2,0.9,0.8045315032435294,1.0,1.2059910039265234
26 health,lin,ERM
   ,2,0.0,0.9,0.8035654031205638,1.0,0.7846793740577468
27 health,lin,ERM
   ,2,0.1,0.9,0.8035654031205638,1.0,0.901715864875564
28 health,lin,ERM
   ,2,0.2,0.9,0.8035654031205638,1.0,1.018752355693381

```

```

29 health,mlp,ERM
30     ,0,0.0,0.996,0.9918441882867648,0.7,0.643148188041 health,lin,AF,1,0.0,9,True,1.0109263503131347
31 health,mlp,ERM
32     ,0,0.1,0.996,0.9918441882867648,0.8,0.846660293642 health,lin,AF,1,0.1,0,True,1.894747702788817
33 health,mlp,ERM
34     ,0,0.2,0.996,0.9918441882867648,0.9,0.884876794248 health,lin,AF,1,0.1,1,True,0.35258218967416993
35 health,mlp,ERM
36     ,1,0.0,0.999,0.9979579146172447,1.0,1.579727959632 health,lin,AF,1,0.1,2,True,1.6004211424178667
37 health,mlp,ERM
38     ,1,0.1,0.999,0.9979579146172447,1.0,1.658783053660 health,lin,AF,1,0.1,3,True,0.3037462982883546
39 health,mlp,ERM
40     ,1,0.2,0.999,0.9979579146172447,1.0,1.658783053668 health,lin,AF,1,0.1,4,True,0.745147641644383
41 health,mlp,ERM
42     ,1,0.3,0.999,0.9979579146172447,1.0,1.658783053670 health,lin,AF,1,0.1,5,True,0.15254389384642275
43 health,mlp,ERM
44     ,1,0.4,0.999,0.9979579146172447,1.0,1.658783053678 health,lin,AF,1,0.1,6,True,1.292150376458172
45 health,mlp,ERM
46     ,1,0.5,0.999,0.9979579146172447,1.0,1.658783053686 health,lin,AF,1,0.1,7,True,0.707239902804329
47 health,mlp,ERM
48     ,1,0.6,0.999,0.9979579146172447,1.0,1.658783053694 health,lin,AF,1,0.1,8,True,1.7981318291546775
49 health,mlp,ERM
50     ,1,0.7,0.999,0.9979579146172447,1.0,1.658783053698 health,lin,AF,1,0.2,0,True,2.0177743179783123
51 health,mlp,ERM
52     ,1,0.8,0.999,0.9979579146172447,1.0,1.658783053702 health,lin,AF,1,0.2,1,True,0.4756088048636653
53 health,mlp,ERM
54     ,1,0.9,0.999,0.9979579146172447,1.0,1.658783053706 health,lin,AF,1,0.2,2,True,1.7234477576073621
55 health,mlp,ERM
56     ,1,0.0,0.999,0.9979579146172447,1.0,1.658783053708 health,lin,AF,1,0.2,3,True,0.42677291347785
57 health,mlp,ERM
58     ,1,0.1,0.999,0.9979579146172447,1.0,1.658783053712 health,lin,AF,1,0.2,4,True,0.8681742568338784
59 health,mlp,ERM
60     ,1,0.2,0.999,0.9979579146172447,1.0,1.658783053716 health,lin,AF,1,0.2,5,True,0.2755705090359181
61 health,mlp,ERM
62     ,1,0.3,0.999,0.9979579146172447,1.0,1.658783053720 health,lin,AF,1,0.2,6,True,1.4151769916476673
63 health,mlp,ERM
64     ,1,0.4,0.999,0.9979579146172447,1.0,1.658783053724 health,lin,AF,1,0.2,7,True,0.8302665179938244
65 health,mlp,ERM
66     ,1,0.5,0.999,0.9979579146172447,1.0,1.658783053728 health,lin,AF,1,0.2,8,True,1.92115844434173
67 health,mlp,ERM
68     ,1,0.6,0.999,0.9979579146172447,1.0,1.658783053732 health,lin,AF,1,0.2,9,True,1.2569795806921256
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117

```

APPENDIX D AUTO-GENERATED INSTANCE COST CSV

Listing 4. Per-instance recourse costs CSV

dataset	model	trainer	seed	epsilon	instance_id	validid	cost
1	health	lin	AF	0,0,0,0	True	0.0022174298096731904	77
2	health	lin	AF	0,0,0,1	True	0.09351049565102389	78
3	health	lin	AF	0,0,0,2	True	0.5008139511917434	79
4	health	lin	AF	0,0,0,3	True	1.7989967006950356	80
5	health	lin	AF	0,0,0,4	True	1.4585106692345036	81
6	health	lin	AF	0,0,0,5	True	0.10900482104156642	82
7	health	lin	AF	0,0,0,6	True	0.009357979114851105	83
8	health	lin	AF	0,0,0,7	True	1.0316063447917123	84
9	health	lin	AF	0,0,0,8	True	0.2709796170170954	85
10	health	lin	AF	0,0,0,9	True	0.1764318752921754	86
11	health	lin	AF	0,0,1,0	True	0.12428426527689317	87
12	health	lin	AF	0,0,1,1	True	0.21557733111824384	88
13	health	lin	AF	0,0,1,2	True	0.6228807866589635	89
14	health	lin	AF	0,0,1,3	True	1.9210635361622554	90
15	health	lin	AF	0,0,1,4	True	1.5805775047017234	91
16	health	lin	AF	0,0,1,5	True	0.23107165650878636	92
17	health	lin	AF	0,0,1,6	True	0.1314248145820711	93
18	health	lin	AF	0,0,1,7	True	1.1536731802589322	94
19	health	lin	AF	0,0,1,8	True	0.3930464524843154	95
20	health	lin	AF	0,0,1,9	True	0.2984987107593954	96
21	health	lin	AF	0,0,2,0	True	0.24635110074411312	97
22	health	lin	AF	0,0,2,1	True	0.33764416658546387	98
23	health	lin	AF	0,0,2,2	True	0.7449476221261834	99
24	health	lin	AF	0,0,2,3	True	2.0431303716294753	100
25	health	lin	AF	0,0,2,4	True	1.7026443401689435	101
26	health	lin	AF	0,0,2,5	True	0.3531384919760064	102
27	health	lin	AF	0,0,2,6	True	0.25349165004929103	103
28	health	lin	AF	0,0,2,7	True	1.2757400157261523	104
29	health	lin	AF	0,0,2,8	True	0.5151132879515354	105
30	health	lin	AF	0,0,2,9	True	0.4205655462266154	106
31	health	lin	AF	0,0,3,0	True	1.7717210875993217	107
32	health	lin	AF	0,0,3,1	True	0.2295555744846745	108
33	health	lin	AF	0,0,3,2	True	1.4773945272283713	109
34	health	lin	AF	0,0,3,3	True	0.1807196830988592	110
35	health	lin	AF	0,0,3,4	True	0.622121026458876	111
36	health	lin	AF	0,0,3,5	True	0.029517278656927354	112
37	health	lin	AF	0,0,3,6	True	1.1691237612686767	113
38	health	lin	AF	0,0,3,7	True	0.5842132876148337	114
39	health	lin	AF	0,0,3,8	True	1.675105213965182	115
40	health	lin	AF	0,0,3,9	True	1.7928175926208496	116
41	health	lin	AF	0,0,4,0	True	1.0109263503131347	117

118	health,mlp,AF,0,0.2,6,True,0.27897143363952637	195	health,lin,ERM,0,0.1,3,True,2.3406677954809383
119	health,mlp,AF,0,0.2,7,False,inf	196	health,lin,ERM,0,0.1,4,True,1.3595109112283565
120	health,mlp,AF,0,0.2,8,False,inf	197	health,lin,ERM,0,0.1,5,True,1.4940954924996346
121	health,mlp,AF,0,0.2,9,True,0.23669147491455078	198	health,lin,ERM,0,0.1,6,True,0.85885855952751913
122	health,mlp,AF,1,0.0,0,True,0.8286165595054626	199	health,lin,ERM,0,0.1,7,True,0.6920265862777982
123	health,mlp,AF,1,0.0,1,True,1.071250319480896	200	health,lin,ERM,0,0.1,8,True,0.16530141673098547
124	health,mlp,AF,1,0.0,2,True,0.6133818626403809	201	health,lin,ERM,0,0.1,9,True,0.6362926644563799
125	health,mlp,AF,1,0.0,3,True,0.0999999403953552	202	health,lin,ERM,0,0.2,0,True,0.9149919730696522
126	health,mlp,AF,1,0.0,4,True,0.38270947337150574	203	health,lin,ERM,0,0.2,1,True,1.079301746466617
127	health,mlp,AF,1,0.0,5,True,0.9619376063346863	204	health,lin,ERM,0,0.2,2,True,1.0248088814723824
128	health,mlp,AF,1,0.0,6,True,2.3279929161071777	205	health,lin,ERM,0,0.2,3,True,2.4547106694724032
129	health,mlp,AF,1,0.0,7,True,0.7694158554077148	206	health,lin,ERM,0,0.2,4,True,1.4735537852198217
130	health,mlp,AF,1,0.0,8,True,0.7187708616256714	207	health,lin,ERM,0,0.2,5,True,1.6081383664910998
131	health,mlp,AF,1,0.0,9,True,0.6485913991928101	208	health,lin,ERM,0,0.2,6,True,0.9729014332666565
132	health,mlp,AF,1,0.1,0,True,0.977965235710144	209	health,lin,ERM,0,0.2,7,True,0.8060694602692634
133	health,mlp,AF,1,0.1,1,True,0.9960500597953796	210	health,lin,ERM,0,0.2,8,True,0.2793442907224506
134	health,mlp,AF,1,0.1,2,True,0.7119928598403931	211	health,lin,ERM,0,0.2,9,True,0.750335538447845
135	health,mlp,AF,1,0.1,3,True,0.1821739375591278	212	health,lin,ERM,1,0.0,0,True,0.4544031954776043
136	health,mlp,AF,1,0.1,4,True,0.49758991599082947	213	health,lin,ERM,1,0.0,1,True,1.1770760395405075
137	health,mlp,AF,1,0.1,5,True,1.1838374137878418	214	health,lin,ERM,1,0.0,2,True,1.5513781335972643
138	health,mlp,AF,1,0.1,6,True,2.5676591396331787	215	health,lin,ERM,1,0.0,3,True,1.3697575946921152
139	health,mlp,AF,1,0.1,7,True,0.8329548835754395	216	health,lin,ERM,1,0.0,4,True,1.677729323247764
140	health,mlp,AF,1,0.1,8,True,0.7808516025543213	217	health,lin,ERM,1,0.0,5,True,0.04324403365116183
141	health,mlp,AF,1,0.1,9,True,0.7513484358787537	218	health,lin,ERM,1,0.0,6,True,0.972477918539937
142	health,mlp,AF,1,0.2,0,True,1.1314444541931152	219	health,lin,ERM,1,0.0,7,True,0.21887793076712805
143	health,mlp,AF,1,0.2,1,True,1.17803955078125	220	health,lin,ERM,1,0.0,8,True,1.2335342986448172
144	health,mlp,AF,1,0.2,2,True,0.7839648127555847	221	health,lin,ERM,1,0.0,9,True,1.02676879701618
145	health,mlp,AF,1,0.2,3,True,0.2460516095161438	222	health,lin,ERM,1,0.1,0,True,0.57113633417937
146	health,mlp,AF,1,0.2,4,True,0.6244319081306458	223	health,lin,ERM,1,0.1,1,True,1.2938091782422734
147	health,mlp,AF,1,0.2,5,True,1.2506043910980225	224	health,lin,ERM,1,0.1,2,True,1.66811127229903
148	health,mlp,AF,1,0.2,6,True,1.1225972175598145	225	health,lin,ERM,1,0.1,3,True,1.4864907333938808
149	health,mlp,AF,1,0.2,7,True,1.010860800743103	226	health,lin,ERM,1,0.1,4,True,1.7944624619495295
150	health,mlp,AF,1,0.2,8,True,0.9001052379608154	227	health,lin,ERM,1,0.1,5,True,0.15997717235292763
151	health,mlp,AF,1,0.2,9,True,0.896099865436554	228	health,lin,ERM,1,0.1,6,True,1.0892110572417029
152	health,mlp,AF,2,0.0,0,True,3.503692388534546	229	health,lin,ERM,1,0.1,7,True,0.33561106946889385
153	health,mlp,AF,2,0.0,1,True,1.1493207216262817	230	health,lin,ERM,1,0.1,8,True,1.3502674373465828
154	health,mlp,AF,2,0.0,2,True,3.8964755535125732	231	health,lin,ERM,1,0.1,9,True,1.1435019357733838
155	health,mlp,AF,2,0.0,3,True,3.741835594177246	232	health,lin,ERM,1,0.2,0,True,0.6878694728811359
156	health,mlp,AF,2,0.0,4,True,3.2063651084899902	233	health,lin,ERM,1,0.2,1,True,1.4105423169440392
157	health,mlp,AF,2,0.0,5,True,2.5007309913635254	234	health,lin,ERM,1,0.2,2,True,1.7848444110007957
158	health,mlp,AF,2,0.0,6,True,5.388555526733398	235	health,lin,ERM,1,0.2,3,True,1.6032238720956469
159	health,mlp,AF,2,0.0,7,True,0.2651277780532837	236	health,lin,ERM,1,0.2,4,True,1.9111956006512953
160	health,mlp,AF,2,0.0,8,True,4.400606155395508	237	health,lin,ERM,1,0.2,5,True,0.2767103110546934
161	health,mlp,AF,2,0.0,9,True,4.526076316833496	238	health,lin,ERM,1,0.2,6,True,1.2059441959434687
162	health,mlp,AF,2,0.1,0,True,4.200447082519531	239	health,lin,ERM,1,0.2,7,True,0.45234420817065973
163	health,mlp,AF,2,0.1,1,True,1.0612828731536865	240	health,lin,ERM,1,0.2,8,True,1.4670005760483487
164	health,mlp,AF,2,0.1,2,True,4.270608425140381	241	health,lin,ERM,1,0.2,9,True,1.2602350744751496
165	health,mlp,AF,2,0.1,3,True,3.9029312133789062	242	health,lin,ERM,2,0.0,0,True,0.056236713989842375
166	health,mlp,AF,2,0.1,4,True,3.3833022117614746	243	health,lin,ERM,2,0.0,1,True,1.7118647040288175
167	health,mlp,AF,2,0.1,5,True,2.5903685092926025	244	health,lin,ERM,2,0.0,2,True,0.4085758944129868
168	health,mlp,AF,2,0.1,6,True,5.484298229217529	245	health,lin,ERM,2,0.0,3,True,0.2698963714817156
169	health,mlp,AF,2,0.1,7,True,0.32856568694114685	246	health,lin,ERM,2,0.0,4,True,0.7548971117045661
170	health,mlp,AF,2,0.1,8,True,4.6868815422058105	247	health,lin,ERM,2,0.0,5,True,0.8574385415387005
171	health,mlp,AF,2,0.1,9,True,4.517419338226318	248	health,lin,ERM,2,0.0,6,True,0.5467518897496327
172	health,mlp,AF,2,0.2,0,True,4.394152641296387	249	health,lin,ERM,2,0.0,7,True,0.665408573976689
173	health,mlp,AF,2,0.2,1,True,1.2017030715942383	250	health,lin,ERM,2,0.0,8,True,0.2177330934524654
174	health,mlp,AF,2,0.2,2,True,4.303710460662842	251	health,lin,ERM,2,0.0,9,True,2.3579908462420525
175	health,mlp,AF,2,0.2,3,True,3.971803665161133	252	health,lin,ERM,2,0.1,0,True,0.17327320480765962
176	health,mlp,AF,2,0.2,4,True,3.635038375854492	253	health,lin,ERM,2,0.1,1,True,1.8289011948466347
177	health,mlp,AF,2,0.2,5,True,2.6649506902071533	254	health,lin,ERM,2,0.1,2,True,0.5256123852308041
178	health,mlp,AF,2,0.2,6,True,5.655788421630859	255	health,lin,ERM,2,0.1,3,True,0.38693286229953283
179	health,mlp,AF,2,0.2,7,True,0.34340187907218933	256	health,lin,ERM,2,0.1,4,True,0.8719336025223833
180	health,mlp,AF,2,0.2,8,True,4.704601287841797	257	health,lin,ERM,2,0.1,5,True,0.9744750323565178
181	health,mlp,AF,2,0.2,9,True,4.650392055511475	258	health,lin,ERM,2,0.1,6,True,0.6637883805674499
182	health,lin,ERM,0,0.0,0,True,0.68690622505867219	259	health,lin,ERM,2,0.1,7,True,0.7824450647945063
183	health,lin,ERM,0,0.0,1,True,0.8512159984836815	260	health,lin,ERM,2,0.1,8,True,0.33476958427028264
184	health,lin,ERM,0,0.0,2,True,0.7967231334894521	261	health,lin,ERM,2,0.1,9,True,2.4750273370598697
185	health,lin,ERM,0,0.0,3,True,2.226624921494973	262	health,lin,ERM,2,0.2,0,True,0.2903096956254768
186	health,lin,ERM,0,0.0,4,True,1.2454680372368914	263	health,lin,ERM,2,0.2,1,True,1.945937685664452
187	health,lin,ERM,0,0.0,5,True,1.3800526185081698	264	health,lin,ERM,2,0.2,2,True,0.6426488760486213
188	health,lin,ERM,0,0.0,6,True,0.74448156852837262	265	health,lin,ERM,2,0.2,3,True,0.5039693531173501
189	health,lin,ERM,0,0.0,7,True,0.5779837122863332	266	health,lin,ERM,2,0.2,4,True,0.9889700933402005
190	health,lin,ERM,0,0.0,8,True,0.05125854273952032	267	health,lin,ERM,2,0.2,5,True,1.091511523174335
191	health,lin,ERM,0,0.0,9,True,0.5222497904649147	268	health,lin,ERM,2,0.2,6,True,0.7808248713852672
192	health,lin,ERM,0,0.1,0,True,0.8009490990781871	269	health,lin,ERM,2,0.2,7,True,0.8994815556123235
193	health,lin,ERM,0,0.1,1,True,0.9652588724751465	270	health,lin,ERM,2,0.2,8,True,0.451806075080809995
194	health,lin,ERM,0,0.1,2,True,0.9107660074809173	271	health,lin,ERM,2,0.2,9,True,2.592063827877687

272 health,mlp, ERM, 0, 0, 0, 0, True, 0.09999999403953552
273 health,mlp, ERM, 0, 0, 0, 1, True, 1.1588484048843384
274 health,mlp, ERM, 0, 0, 0, 2, False, inf
275 health,mlp, ERM, 0, 0, 0, 3, False, inf
276 health,mlp, ERM, 0, 0, 0, 4, True, 0.09999999403953552
277 health,mlp, ERM, 0, 0, 0, 5, False, inf
278 health,mlp, ERM, 0, 0, 0, 6, True, 1.7681208848953247
279 health,mlp, ERM, 0, 0, 0, 7, True, 0.1756606101989746
280 health,mlp, ERM, 0, 0, 0, 8, True, 0.09999999403953552
281 health,mlp, ERM, 0, 0, 0, 9, True, 1.099407434463501
282 health,mlp, ERM, 0, 0, 1, 0, True, 0.2847887873649597
283 health,mlp, ERM, 0, 0, 1, 1, True, 1.350301742553711
284 health,mlp, ERM, 0, 0, 1, 2, False, inf
285 health,mlp, ERM, 0, 0, 1, 3, False, inf
286 health,mlp, ERM, 0, 0, 1, 4, True, 0.17709362506866455
287 health,mlp, ERM, 0, 0, 1, 5, True, 1.4368488788604736
288 health,mlp, ERM, 0, 0, 1, 6, True, 1.783553123474121
289 health,mlp, ERM, 0, 0, 1, 7, True, 0.3221847414970398
290 health,mlp, ERM, 0, 0, 1, 8, True, 0.14391285181045532
291 health,mlp, ERM, 0, 0, 1, 9, True, 1.2745985984802246
292 health,mlp, ERM, 0, 0, 2, 0, True, 0.2927534878253937
293 health,mlp, ERM, 0, 0, 2, 1, True, 1.420407772064209
294 health,mlp, ERM, 0, 0, 2, 2, False, inf
295 health,mlp, ERM, 0, 0, 2, 3, True, 1.6530849933624268
296 health,mlp, ERM, 0, 0, 2, 4, True, 0.37585678696632385
297 health,mlp, ERM, 0, 0, 2, 5, True, 1.0841124057769775
298 health,mlp, ERM, 0, 0, 2, 6, True, 2.038994312286377
299 health,mlp, ERM, 0, 0, 2, 7, True, 0.4549943208694458
300 health,mlp, ERM, 0, 0, 2, 8, True, 0.09999996423721313
301 health,mlp, ERM, 0, 0, 2, 9, True, 0.543687105178833
302 health,mlp, ERM, 1, 0, 0, 0, True, 0.6265113949775696
303 health,mlp, ERM, 1, 0, 0, 1, True, 0.5598059892654419
304 health,mlp, ERM, 1, 0, 0, 2, True, 4.995126724243164
305 health,mlp, ERM, 1, 0, 0, 3, True, 0.8848893046379089
306 health,mlp, ERM, 1, 0, 0, 4, True, 0.9069631099700928
307 health,mlp, ERM, 1, 0, 0, 5, True, 6.042341232299805
308 health,mlp, ERM, 1, 0, 0, 6, True, 0.9130632877349854
309 health,mlp, ERM, 1, 0, 0, 7, True, 0.42502132058143616
310 health,mlp, ERM, 1, 0, 0, 8, True, 0.09999999403953552
311 health,mlp, ERM, 1, 0, 0, 9, True, 0.3435572385787964
312 health,mlp, ERM, 1, 0, 1, 0, True, 0.7597363591914153
313 health,mlp, ERM, 1, 0, 1, 1, True, 0.7501630187034607
314 health,mlp, ERM, 1, 0, 1, 2, True, 4.815957069396973
315 health,mlp, ERM, 1, 0, 1, 3, True, 1.1189303398132324
316 health,mlp, ERM, 1, 0, 1, 4, True, 0.8854197263717651
317 health,mlp, ERM, 1, 0, 1, 5, True, 6.365862846374512
318 health,mlp, ERM, 1, 0, 1, 6, True, 0.88718813657776062
319 health,mlp, ERM, 1, 0, 1, 7, True, 0.5466513633728027
320 health,mlp, ERM, 1, 0, 1, 8, True, 0.09999998658895493
321 health,mlp, ERM, 1, 0, 1, 9, True, 0.35792168974876404
322 health,mlp, ERM, 1, 0, 2, 0, True, 0.8752454519271851
323 health,mlp, ERM, 1, 0, 2, 1, True, 0.7518577575683594
324 health,mlp, ERM, 1, 0, 2, 2, True, 4.760556697845459
325 health,mlp, ERM, 1, 0, 2, 3, True, 1.0518792867660522
326 health,mlp, ERM, 1, 0, 2, 4, True, 0.9736133813858032
327 health,mlp, ERM, 1, 0, 2, 5, True, 2.644326686859131
328 health,mlp, ERM, 1, 0, 2, 6, True, 0.7649006843566895
329 health,mlp, ERM, 1, 0, 2, 7, True, 0.6719211935997009
330 health,mlp, ERM, 1, 0, 2, 8, True, 0.09999998658895493
331 health,mlp, ERM, 1, 0, 2, 9, True, 0.40259477496147156
332 health,mlp, ERM, 2, 0, 0, 0, True, 1.6526018381118774
333 health,mlp, ERM, 2, 0, 0, 1, True, 1.0587971210479736
334 health,mlp, ERM, 2, 0, 0, 2, True, 0.871553897857666
335 health,mlp, ERM, 2, 0, 0, 3, True, 1.5454106330871582
336 health,mlp, ERM, 2, 0, 0, 4, True, 0.09999999403953552
337 health,mlp, ERM, 2, 0, 0, 5, True, 4.139216899871826
338 health,mlp, ERM, 2, 0, 0, 6, True, 1.2884514331817627
339 health,mlp, ERM, 2, 0, 0, 7, True, 1.023694634437561
340 health,mlp, ERM, 2, 0, 0, 8, True, 0.09999997913837433
341 health,mlp, ERM, 2, 0, 0, 9, False, inf
342 health,mlp, ERM, 2, 0, 1, 0, True, 2.3038556575775146
343 health,mlp, ERM, 2, 0, 1, 1, True, 0.48094290494918823
344 health,mlp, ERM, 2, 0, 1, 2, True, 0.9778098464012146
345 health,mlp, ERM, 2, 0, 1, 3, True, 1.8084425926208496
346 health,mlp, ERM, 2, 0, 1, 4, True, 0.19984257221221924
347 health,mlp, ERM, 2, 0, 1, 5, True, 4.7877950668333496
348 health,mlp, ERM, 2, 0, 1, 6, True, 1.3899121284484863

```
349 | health,mlp,ERM,2,0.1,7,True,1.2289860248565674  
350 | health,mlp,ERM,2,0.1,8,True,0.291856586933136  
351 | health,mlp,ERM,2,0.1,9,False,inf  
352 | health,mlp,ERM,2,0.2,0,False,inf  
353 | health,mlp,ERM,2,0.2,1,True,0.4565989375114441  
354 | health,mlp,ERM,2,0.2,2,True,1.0854618549346924  
355 | health,mlp,ERM,2,0.2,3,True,1.8083579540252686  
356 | health,mlp,ERM,2,0.2,4,False,inf  
357 | health,mlp,ERM,2,0.2,5,True,1.6818499565124512  
358 | health,mlp,ERM,2,0.2,6,True,1.5481600761413574  
359 | health,mlp,ERM,2,0.2,7,True,1.1980092525482178  
360 | health,mlp,ERM,2,0.2,8,True,0.37811192870140076  
361 | health,mlp,ERM,2,0.2,9,True,1.961012840270996
```

APPENDIX E

AUTO-GENERATED ACTION PROFILE CSV

Listing 5. Feature-wise action diagnostics CSV

ACKNOWLEDGMENT

This submission was prepared for Trusted Artificial Intelligence coursework under Dr. Mostafa Tavasolipour.

REFERENCES

- [1] B. Ustun, A. Spangher, and Y. Liu, "Actionable recourse in linear classification," in *Proceedings of the Conference on Fairness, Accountability, and Transparency*, 2019.
 - [2] J. Pearl, *Causality: Models, Reasoning, and Inference*. Cambridge University Press, 2009.
 - [3] A.-H. Karimi, G. Barthe, B. Balle, and I. Valera, "Algorithmic recourse: from counterfactual explanations to interventions," in *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, 2021.
 - [4] R. K. Mothilal, A. Sharma, and C. Tan, "Explaining machine learning classifiers through diverse counterfactual explanations," *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency*, 2020.
 - [5] A.-H. Karimi, B. Schölkopf, and I. Valera, "Algorithmic recourse under imperfect causal knowledge: A probabilistic approach," in *Advances in Neural Information Processing Systems*, 2020.